

# **CS354**

# **Operating Systems**

**<http://www.cs.purdue.edu/people/comer>**

# Module I

## Overview

# Topics

- Introductions
- Course motivation and scope
- Xinu and the lab
- Required background
- The drop/stay decision
- Syllabus

# INTRODUCTIONS

# **COURSE MOTIVATION AND SCOPE**

# Scope

This is a course about the design and structure of computer operating systems. It covers the concepts, principles, functionality, tradeoffs, and implementation of systems that support concurrent processing.

# What We Will Cover

- Operating system design
- Functionality an operating system offers
- Major system components
- Interdependencies and system structure
- The key relationships between operating system abstractions and the underlying hardware (especially processes and interrupts)
- Implementation details

# What You Will Learn

- Fundamental
  - Principles
  - Design options
  - Tradeoffs
- How to modify and test operating system code
- How to design and build an operating system



# What We Will NOT Cover

- The course is not
  - A comparison of large commercial and open source operating systems
  - A description of features or instructions on how to use a particular operating system
  - A survey of research systems and alternative approaches that have been studied
  - A set of techniques for building operating systems on unusual hardware

# How Did We Get Here?

# How Did We Get Here?

- 1940s – 1950s: How can computers be built?
  - The dawn of digital computers
  - Each processor special-purpose
  - Unique interface for each I/O device
  - Software written in assembly language
  - Programs written to control the hardware

# How Did We Get Here?

## (continued)

- 1960s: What are the best abstractions to use?
  - A move to general-purpose hardware
  - *Instruction Set Architecture* emerges
  - Families of computers devised
  - I/O became independent of specific hardware details
  - High-level programming languages created (e.g., parameterized functions, data types, and recursion)

# How Did We Get Here?

## (continued)

- Late 1960s and 1970s
  - Researchers investigate operating systems
  - Multics project devises brilliant abstractions
  - Unix introduces simple, elegant, and efficient versions of the Multics abstractions
- 1980s and on
  - The Internet project discovers a set of communication abstractions that are elegant and efficient
  - All systems are connected

# What Can We Conclude From History?

# What Can We Conclude From History?

- The gap between hardware and users is huge

# What Can We Conclude From History?

- The gap between hardware and users is huge
  - Hardware is ugly



# What Can We Conclude From History?

- The gap between hardware and users is huge
  - Hardware is ugly
  - Operating system abstractions are beautiful

# What Can We Conclude From History?

- The gap between hardware and users is huge
  - Hardware is ugly
  - Operating system abstractions are beautiful
- Everything is connected

# What Can We Conclude From History?

- The gap between hardware and users is huge
  - Hardware is ugly
  - Operating system abstractions are beautiful
- Everything is connected
  - Data centers

# What Can We Conclude From History?

- The gap between hardware and users is huge
  - Hardware is ugly
  - Operating system abstractions are beautiful
- Everything is connected
  - Data centers
  - Desktops

# What Can We Conclude From History?

- The gap between hardware and users is huge
  - Hardware is ugly
  - Operating system abstractions are beautiful
- Everything is connected
  - Data centers
  - Desktops
  - Smart phones

# What Can We Conclude From History?

- The gap between hardware and users is huge
  - Hardware is ugly
  - Operating system abstractions are beautiful
- Everything is connected
  - Data centers
  - Desktops
  - Smart phones
  - Devices

# What Can We Conclude From History?

- The gap between hardware and users is huge
  - Hardware is ugly
  - Operating system abstractions are beautiful
- Everything is connected
  - Data centers
  - Desktops
  - Smart phones
  - Devices
  - Sensors

# Consequences And Observations



# Consequences And Observations

- Our job in Computer Science is to build beautiful new abstractions that fill the gap between ugly hardware and users

# Consequences And Observations

- Our job in Computer Science is to build beautiful new abstractions that fill the gap between ugly hardware and users
- We cannot consider operating systems without including computer networking

# Consequences And Observations

- Our job in Computer Science is to build beautiful new abstractions that fill the gap between ugly hardware and users
- We cannot consider operating systems without including computer networking
- The central questions in computing have always focused on the tradeoff between imagined beauty and performance

# Consequences And Observations

- Our job in Computer Science is to build beautiful new abstractions that fill the gap between ugly hardware and users
- We cannot consider operating systems without including computer networking
- The central questions in computing have always focused on the tradeoff between imagined beauty and performance
  - It is easy to imagine new abstractions

# Consequences And Observations

- Our job in Computer Science is to build beautiful new abstractions that fill the gap between ugly hardware and users
- We cannot consider operating systems without including computer networking
- The central questions in computing have always focused on the tradeoff between imagined beauty and performance
  - It is easy to imagine new abstractions
  - We must restrict ourselves to abstractions that map onto the underlying hardware efficiently (and hope that hardware engineers eventually adapt to our abstractions)

# The Once And Future Hot Topic

- In the 1970s and early 1980s, operating systems was one of the hottest topics in CS
- By the mid-1990s, OS research had stagnated
- Now things have heated up again, and new operating systems are being designed for
  - Multicore systems
  - Data centers
  - Smart phones
  - Large and small embedded devices
  - The Internet of Things

# **XINU AND THE LAB**

# Motivation For Using A Real System

- Provides examples of the principles
- Makes everything concrete
- Gives students a chance to experiment
- Shows how abstractions map to real hardware



# Why Xinu?

- Small — can be read and understood in a semester
- Complete — includes all the major components
- Elegant — provides an excellent example of clean design
- Powerful — has dynamic process creation, dynamic memory management, and flexible I/O
- Practical — has been used in real products

# The Xinu Lab

- Innovative facility for rapid OS development and testing
- Completely automated
- Allows each student to create, download, and run code on bare hardware
- Handles hardware reboot when necessary
- Includes connections to the Internet

# **REQUIRED BACKGROUND AND PREREQUISITES**

# Background Needed

- Concepts from undergrad operating systems
  - Concurrent programming: you should have written a program that uses *fork* or the equivalent
  - Understanding of deadlock and race conditions
  - I/O: you should know the difference between standard library functions (e.g., *fopen*, *putc*, *getc*, *fread*, *fwrite*) and system calls (e.g., *open*, *close*, *read*, *write*)
  - File systems and hierarchical directories
  - Symbolic and hard links
  - File modes and protection

# Background Needed (continued)

- Understanding of runtime storage components
  - Segments (text, data, and bss) and their layout
  - Runtime stack, activation records, and argument passing
  - Basic heap storage management (free list)
- C programming
  - At least one nontrivial program
  - Comfortable with low-level constructs (e.g., bit manipulation and pointers)

# Background Needed (continued)

- Working knowledge of basic UNIX tools
  - Text editor (e.g., emacs)
  - Compiler / linker / loader
  - Tar archives
  - Make and Makefiles
- Desire to learn

# **SYLLABUS**

# How We Will Proceed

- We will examine the major components of an operating system
- For a given component we will
  - Outline the functionality needed
  - Learn the key principles involved
  - Understand one particular design choice in depth
  - Consider implementation details and the relationship to hardware
  - Discuss other possibilities and tradeoffs
- Note: we will cover components in a linear order that allows us to understand one component at a time without relying on later components





**Questions?**

**A FEW THINGS  
TO THINK ABOUT**

**Perfection [in design] is achieved not when there is nothing to add,  
but rather when there is nothing more to take away.**

*– Antoine de Saint-Exupery*

**A teacher's job is to make the agony of decision making so intense  
you can only escape by thinking.**

*– source unknown*

**Real concurrency — in which one program actually continues to function while you call up and use another — is more amazing but of small use to the average person. How many programs do you have that take more than a few seconds to perform any task?**

*(From an article about new operating systems for the IBM PC in the New York Times, 25 April 1989)*

# **Module II**

## **Organization Of An Operating System**

# What Is An Operating System?

- Provides abstract computing environment
- Supplies computational services
- Manages resources
- Hides low-level hardware details
- Note: operating system software is among the most complex ever devised

# Example Services An OS Supplies

- Support for concurrent execution
- Process synchronization
- Inter-process communication mechanisms
- Message passing and asynchronous events
- Management of address spaces and virtual memory
- Protection among users and running applications
- High-level interface for I/O devices
- A file system and file access facilities
- Intermachine communication



# What An Operating System Is NOT

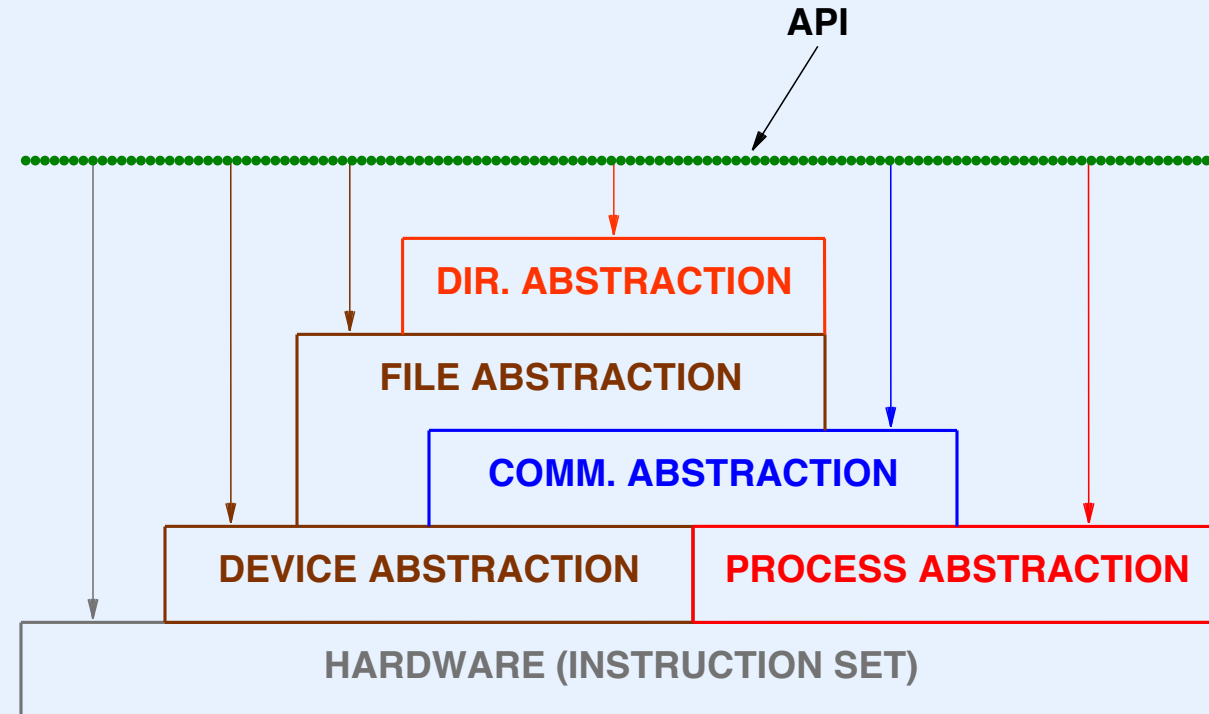
- Hardware
- Language
- Compiler
- Windowing system or browser
- Command interpreter
- Library of utility functions
- Graphical desktop

# **AN OPERATING SYSTEM FROM THE OUTSIDE**

# The System Interface

- Single copy of OS per computer
  - Hidden from users
  - Accessible only to application programs
- *Application Program Interface (API)*
  - Defines services OS makes available
  - Defines parameters for the services
  - Provides access to all abstractions
  - Hides hardware details

# OS Abstractions And Application Interface



- Modules in the OS offer services
- Some services build on others

# Interface To System Services

- Appears to operate like a function call mechanism
  - OS makes set of “functions” available to applications
  - Application supplies arguments using standard mechanism
  - Application “calls” one of the OS functions
- Control transfers to OS code that implements the function
- Control returns to caller when function completes

# Interface To System Services (continued)

- Requires special instruction to invoke OS function
  - Moves from application *address space* to OS
  - Changes from application *mode* or *privilege level* to OS
- Terminology used by various vendors
  - *System call*
  - *Trap*
  - *Supervisor call*
- We will use the generic term *system call*

# Example System Call In Xinu: Write A Character On The Console

```
/* ex1.c - main */  
  
#include <xinu.h>  
  
/*-----  
 * main - Write "hi" on the console  
 *-----  
 */  
void main(void)  
{  
    putc(CONSOLE, 'h');  
    putc(CONSOLE, 'i');  
    putc(CONSOLE, '\n');  
}
```

- Note: we will discuss the implementation of *putc* later.

# OS Services And System Calls

- Each OS service accessed through system call interface
- Most services employ a set of several system calls
- Examples
  - Process management service includes functions to *suspend* and then *resume* a process
  - *Socket API* used for Internet communication includes many functions



# System Calls Used With I/O

- Open-close-read-write paradigm
- Application
  - Uses *open* to connect to a file or device
  - Calls functions to *write* data or *read* data
  - Calls *close* to terminate use
- Internally, the set of I/O functions coordinate
  - *Open* returns a descriptor, *d*
  - *Read* and *write* operate on descriptor *d*

# Questions

- How many system calls does an OS need?
- What should they be?

# Concurrent Processing

- Fundamental concept that dominates OS design
- *Real concurrency* achieved by hardware
  - I/O devices operate at same time as processor
  - Multiple processors/cores each operate at the same time
- *Apparent concurrency* achieved with multitasking (multiprogramming)
  - Multiple programs appear to operate simultaneously
  - Operating system provides the illusion

# Multitasking

- Powerful abstraction
- Allows user(s) to run multiple computations
- OS switches processor(s) among available computations quickly
- All computations appear to proceed in parallel

# Terminology Used With Multitasking

- A *program* consists of static code and data
- A *function* is a unit of application program code
- A *process* (also called a *thread of execution*) is an active computation (i.e., the execution or “running” of a program)

# Process

- OS abstraction
- Created by OS system call
- Managed entirely by OS; unknown to hardware
- Operates concurrently with other processes

# Example Of Process Creation In Xinu (Part 1)

```
/* ex2.c - main, sndA, sndB */

#include <xinu.h>

void    sndA(void), sndB(void);

/*-----
 * main - Example of creating processes in Xinu
 *-----
 */
void    main(void)
{
    resume( create(sndA, 1024, 20, "process 1", 0) );
    resume( create(sndB, 1024, 20, "process 2", 0) );
}

/*-----
 * sndA - Repeatedly emit 'A' on the console without terminating
 *-----
 */
void    sndA(void)
{
    while( 1 )
        putc(CONSOLE, 'A');
}
```

# Example Of Process Creation In Xinu (Part 2)

```
/*-----  
 * sndB - Repeatedly emit 'B' on the console without terminating  
 *-----  
 */  
void    sndB(void)  
{  
    while( 1 )  
        putc(CONSOLE, 'B');  
}
```



# Difference Between Function Call And Process Creation

- Normal function call
  - Synchronous execution
  - Single computation
- The *create* system call that starts a new process
  - Asynchronous execution
  - Two processes proceed after the call

# Distinction Between A Program And A Process

- Sequential program
  - Set of functions executed by a single thread of control
- Process
  - Computational abstraction not usually part of the programming language
  - Created independent of code that is executed
  - Key idea: multiple processes can execute the same code concurrently
- In the following example, two processes execute function *sndch* concurrently

# Example Of Two Processes Sharing Code

```
/* ex3.c - main, sndch */

#include <xinu.h>

void    sndch(char);

/*-----
 * main    -   Example of 2 processes executing the same code concurrently
 *-----
 */
void    main(void)
{
    resume( create(sndch, 1024, 20, "send A", 1, 'A') );
    resume( create(sndch, 1024, 20, "send B", 1, 'B') );
}

/*-----
 * sndch    -   Output a character on a serial device indefinitely
 *-----
 */
void    sndch(
    char    ch                /* The character to emit continuously */
)
{
    while ( 1 )
        putc(CONSOLE, ch);
}
```

# Storage Allocation When Multiple Processes Execute

- Various memory models exist for multitasking environments
- Each process requires its own
  - Runtime stack for function calls
  - Storage for local variables
  - Copy of arguments
- A process *may* have private heap storage as well

# Consequence For Programmers

A copy of function arguments and local variables are associated with each process executing a particular function, *not* with the code in which they are declared.

# **AN OPERATING SYSTEM FROM THE INSIDE**

# Operating System

- Well-understood subsystems
- Many subsystems employ heuristic policies
  - Policies can conflict
  - Heuristics can have corner cases
- Complexity arises from interactions among subsystems
- Side-effects can be
  - Unintended
  - Unanticipated
- We will see examples

# Building An Operating System



# Building An Operating System

- The intellectual challenge comes from the design of a “system” rather than from the design individual pieces
- Structured design is needed
- It can be difficult to understand the consequences of choices
- We will use a hierarchical microkernel design to help control complexity

# Major OS Components

- Process manager
- Memory manager
- Device manger
- Clock (time) manager
- File manager
- Interprocess communication
- Intermachine communication
- Assessment and accounting

# Multilevel Structure

- Organizes components
- Controls interactions among subsystems
- Allows a system to be understood and built incrementally
- Differs from traditional layered approach
- Will be employed as the design paradigm throughout the text and course

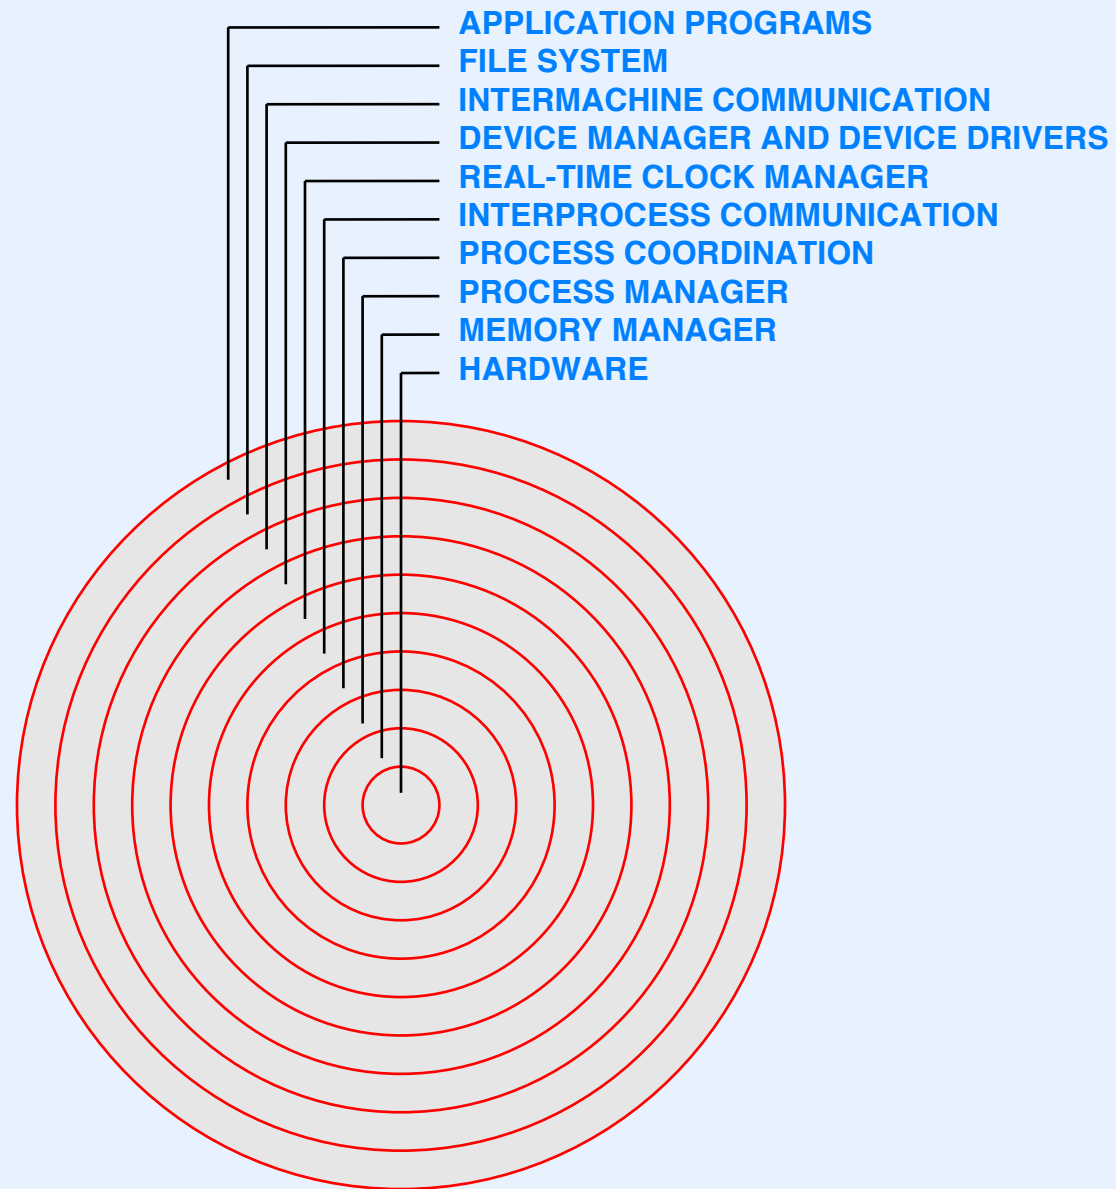
# Multilevel Vs. Multilayered Organization

- Multilayer structure
  - Visible to user as well as designer
  - Each layer uses layer directly beneath
  - Involves protection as well as data abstraction
  - Examples
    - \* Internet protocol layering
    - \* MULTICS layered security structure
  - Can be inefficient

# Multilevel Vs. Multilayered Organization (continued)

- Multilevel structure
  - Form of data abstraction
  - Used during system construction
  - Helps designer focus attention on one aspect at a time
  - Keeps policy decisions independent
  - Allows given level to use *all* lower levels
  - Efficient

# Multilevel Structure Of Xinu



# How To Build An OS

- Work one level at a time
- Identify a service to be provided
- Begin with a *philosophy*
- Establish *policies* that follow the philosophy
- Design *mechanisms* that enforce the policies
- Construct an *implementation* for specific hardware

# Design Example

- Example: access to I/O
- Philosophy: “fairness”
- Policy: FCFS resource access
- Mechanism: queue of requests (FIFO)
- Implementation: program written in C



# Summary

- Operating system supplies set of services
- System calls provide interface between OS and application
- Concurrency is fundamental concept
  - Between I/O devices and processor
  - Between multiple computations
- Process is OS abstraction for concurrency
- Process differs from program or function
- You will learn how to design and implement system software that supports concurrent processing


## Summary (continued)

- OS has well-understood internal components
- Complexity arises from interactions among components
- Multilevel approach helps organize system structure
- Design involves inventing policies and mechanisms that enforce overall goals
- Xinu includes a compact list structure that uses relative pointers and an implicit data structure to reduce size
- Xinu type names specify both purpose and data size

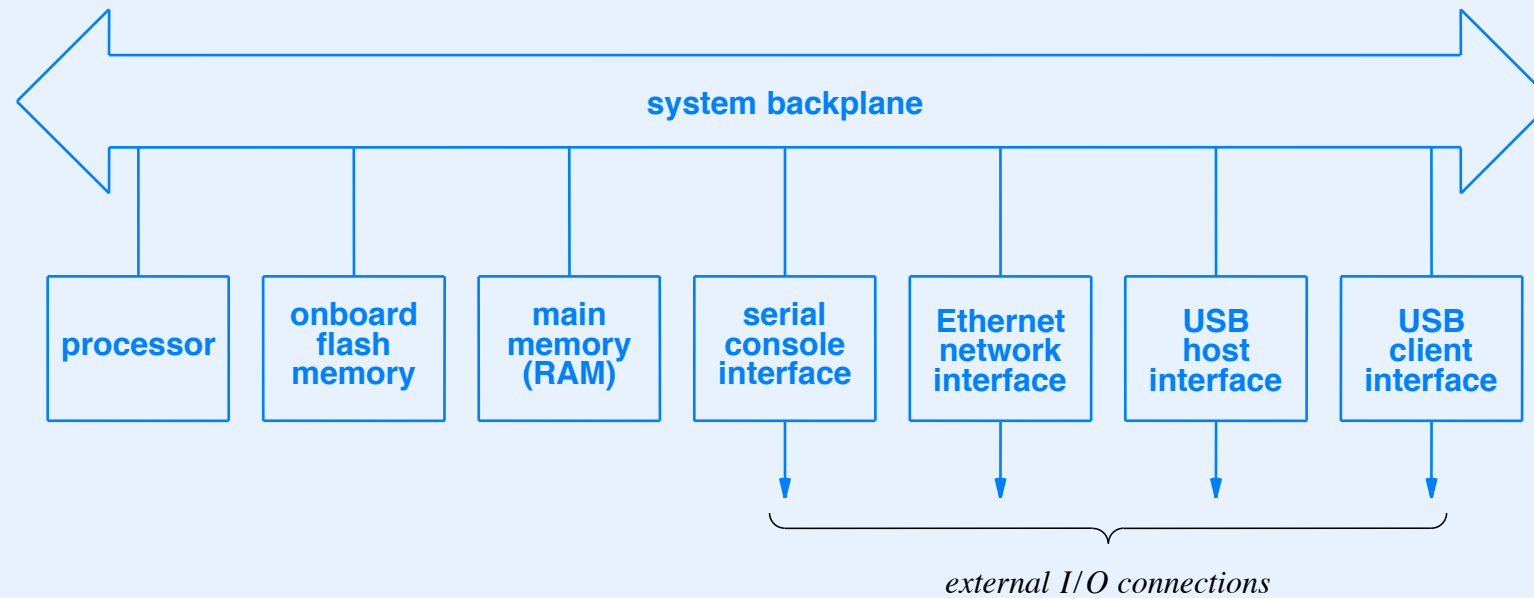
# **Module III**

## **Hardware Architecture And Runtime Systems (A Brief Overview)**

# Location Of Hardware In The Hierarchy



# Organization Of Major Hardware Components



- Example platforms use *System on Chip (SoC)* approach
- Single VLSI chip contains components and interconnections
- Not relevant to operating system

# Key Hardware Features From An Operating System Point Of View

- Processor
- Memory system
- Bus and address space
- Individual I/O devices
- Interrupt structure

# Processor

- Instruction set
- Registers
- Optional hardware facilities
  - Firmware (e.g., BIOS in ROM)
  - Co-processor
  - Interrupt processor

# Instruction Sets On The Example Platforms

- Galileo board
  - CISC design
  - Well-known Intel (x86) instruction set
- BeagleBone Black
  - RISC design
  - Well-known ARM instruction set
- Note: we will see how low-level operating system functions differ on the two



# General-purpose And Special-purpose Registers

- General-purpose registers
  - Local storage inside processor
  - Hold active values during computation (e.g., used to compute an expression)
  - Saved and restored during subprogram invocation
- Special-purpose registers
  - Located inside the processor
  - Values control processor actions (e.g., mode and address space visibility)

# Memory System

- Defines size of a *byte*, the smallest addressable unit
- Provides address space
- Typical physical address space is
  - *Monolithic*
  - *Linear* (but may not be contiguous)
- Includes caching
- Defines important property for programmers: endianness


# Byte Order Terminology

- Order of bytes within an integer in memory
- Irrelevant to data in registers
- *Little Endian* stores least-significant byte at lowest address
- *Big Endian* stores most-significant byte at lowest address

# Memory Caches

- Special-purpose hardware units
- Speed memory access
- Less expensive than high-speed memory
- Physically associated with memory
- Conceptually placed “between” processor and memory

# Conceptual Placement Of Memory Cache



- All references (including instruction fetch) go through cache
- Multi-level cache possible
- In practice, placing the cache on the processor chip improves performance
- Key question: are virtual or physical addresses cached?
- Consequence: operating system may need to
  - Flush the cache
  - Avoid the cache when accessing a device

# I/O Devices

- Wide variety of peripheral devices available
  - Displays
  - Keyboards / mice
  - Disks
  - Wired and wireless network interfaces
  - Printers and scanners
  - Cameras
  - Sensors
- Multiple transfer paradigms (character, block, packet, stream)

# Communication Between Device And Processor

- Communication with I/O devices is *memory-mapped*
- Processor can
  - Interrogate device status
  - Control (e.g., shutdown or awaken) a device
  - Start a data transfer
- Device uses bus to
  - Reply to commands from the processor
  - Transfer data to/ from memory
  - Interrupt when operation completes

# Bus Operations

- Only two basic operations supported
  - Fetch
  - Store
- All I/O uses fetch-store paradigm
- Fetch
  - Processor places address on bus
  - Processor uses control line to signal *fetch request*
  - Device senses its address
  - Device puts specified data on bus
  - Device uses control line to signal *response*



# Bus Operations

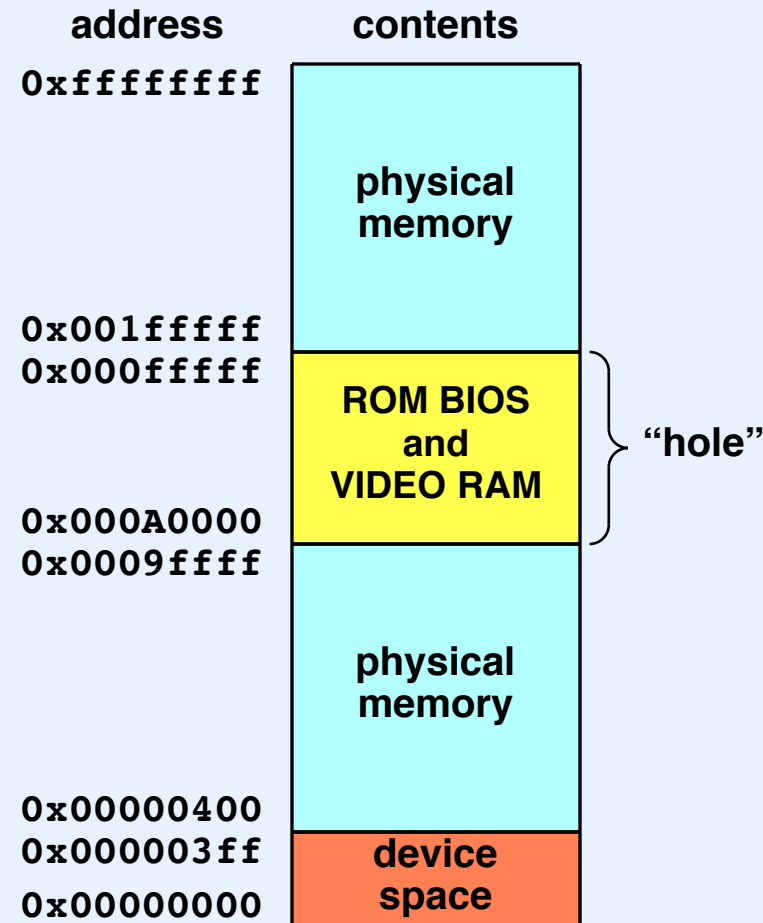
## (continued)

- Store
  - Processor places data and address on bus
  - Processor uses control line to signal *store request*
  - Device senses its address
  - Device extracts data from the bus
  - Device uses control line to signal *data extracted*

# Two basic Approaches For I/O

- *Port-mapped I/O*
  - Special instructions used to access devices
  - Used on earlier Intel machines
- *Memory-mapped I/O*
  - Devices placed beyond physical memory
  - Processor uses normal fetch / store memory instructions
  - On later Intel machines as well as most others

# Address Space On An Intel System

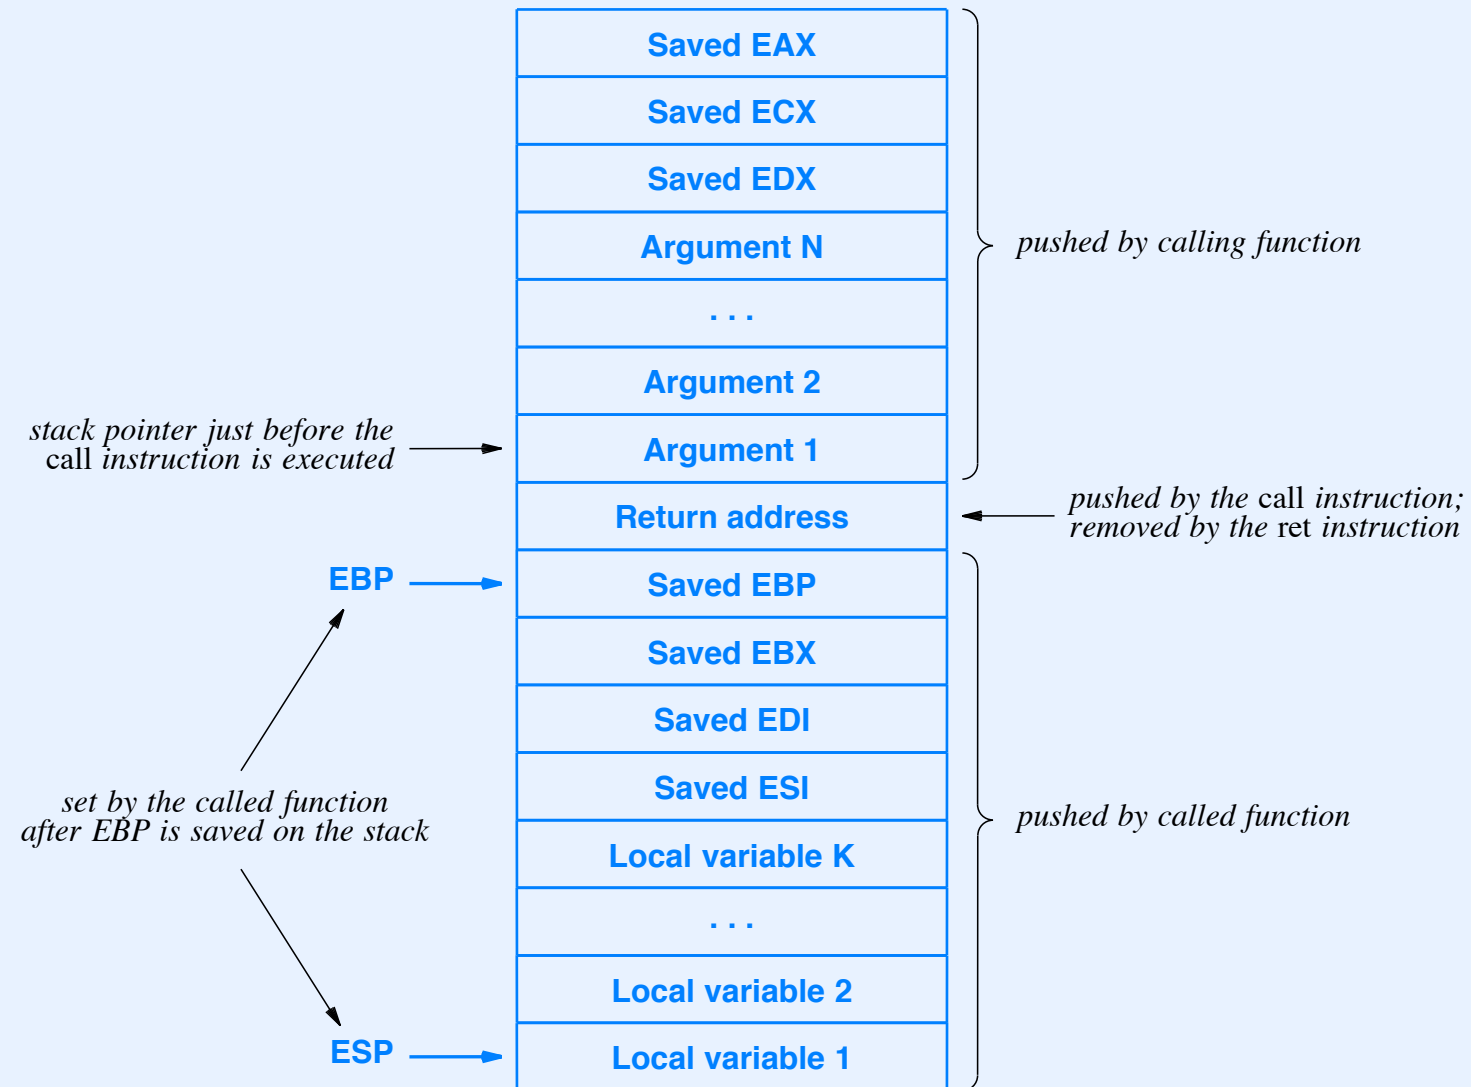


- Often discontinuous
- Traditional PCs have a “hole” from 640KB to 1MB

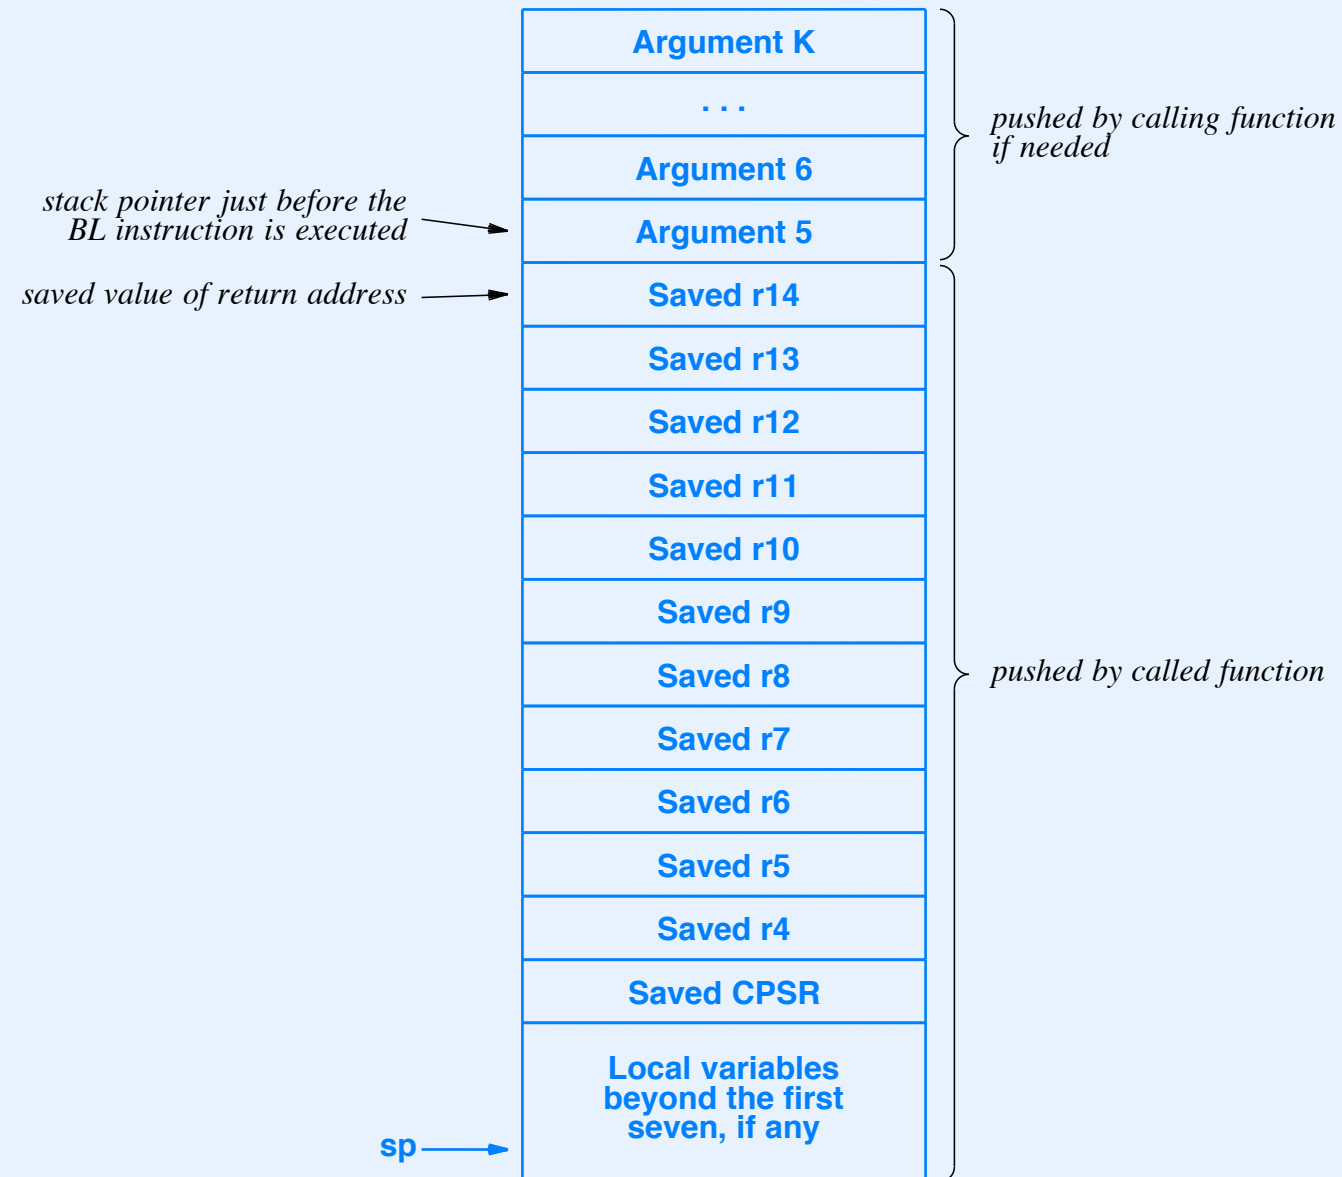
# Calling Conventions

- Refer to the set of steps taken during a function call
- The conventions specify
  - Which registers must be saved by caller
  - Which registers must be saved by called function
  - Where registers are saved (e.g., on the stack)
- Specific details may depend on
  - The hardware
  - The compiler being used
- We will see that key operating systems functions cannot be implemented unless the calling conventions are known

# Calling Conventions For Intel



# Calling Conventions For ARM



# Interrupt Mechanism

- Fundamental role in modern system
- Permits I / O concurrent with execution
- Allows device priority
- Informs processor when I / O finished
- *Software interrupt* also possible

# Interrupt-Driven I/O

- Processor starts device
- Device operates concurrently
- Device interrupts the processor when finished with the assigned task
- Interrupt timing
  - Asynchronous wrt computation
  - Synchronous wrt an individual instruction (occurs between instructions)



# Interrupt Mask

- Bit mask kept in a processor status register
- Determines whether interrupts are enabled
- Hardware sets mask during interrupt to prevent further interrupts
- Interrupt priority scheme
  - Offered by some hardware
  - Each device assigned priority level (binary number)
  - When servicing level  $K$  interrupt, hardware sets mask to disable interrupts at level  $K$  and lower

# Interrupt Processing

- Operating system must
  - Store address of interrupt code in *vector* for each device
  - Arrange for interrupt code to save registers used during the interrupt and restore registers before returning
- More details later in the course

# Returning From An Interrupt

- Special hardware instruction used
- Atomically restores
  - Old program state
  - Interrupt mask
  - Instruction pointer (program counter)
- Hardware returns to location where interrupt occurred, and processing continues exactly as if no interrupt happened

# Transfer Size And Interrupts

- Interrupt occurs after I/O operation completes
- Transfer size depends on device
  - Serial device transfers one character
  - Disk controller transfers one block (e.g., 512 bytes)
  - Network interface transfers one packet
- Large transfers use *Direct Memory Access (DMA)*


# Direct Memory Access (DMA)

- Hardware mechanism
- I/O device transfers large block of data to / from memory without using the processor
- Example: network interface places incoming network packet in memory buffer
- Advantage: allows processor to execute during I/O transfer
- Disadvantages
  - More expensive
  - More complex to design and program

# Memory Segments In C Programs

- C Program has four primary data areas called *segments*
- Text segment
  - Contains program code
  - Usually unwritable
- Data segment
  - Contains initialized data values (globals)
- Bss segment
  - Contains uninitialized data values (set to zero)
- Stack segment
  - Used for function calls

# Storage Layout For A C Program



- Stack grows downward (toward lower memory addresses)
- Heap grows upward

# Symbols For Segment Addresses

- C compiler and / or linker adds three reserved names to symbol table
- `_etext` lies beyond text segment
- `_edata` lies beyond data segment
- `_end` lies beyond bss segment
- Only the addresses are significant; values are irrelevant
- Program can use the addresses of the reserved symbol to determine the size of segments
- Note: names are declared to be *extern* without the underscore:


```
extern int end;
```



# Runtime Storage For A Process

- Text is shared
- Stack cannot be shared
- Data area *may* be shared
- Exact details depend on address space model OS offers

# Example Runtime Storage Model: Xinu



- Single, shared copy of
  - Text segment
  - Data segment
  - Bss segment
- One stack segment per process
  - Allocated when process created
  - Each process has its own stack pointer

# Summary

- Components of third generation computer
  - Processor
  - Main memory
  - I / O Devices
    - \* Accessed over bus
    - \* Operate concurrently with processor
    - \* Can be memory-mapped or port-mapped
    - \* Can use DMA
    - \* Employ interrupts

# Summary

## (continued)

- Interrupt mechanism
  - Informs processor when I / O completes
  - Permits asynchronous device operation
- C uses four memory areas: text, data, bss, and stack segments
- Multiple concurrent computations
  - Can share text, data, and bss
  - Cannot share stack



**Questions?**