


# **Module V**

## **Process Coordination And Synchronization**

# Location Of Process Coordination In The Hierarchy



# Coordination Of Processes

- Necessary in a concurrent system
- Avoids conflicts when accessing shared items
- Allows multiple processes to cooperate
- Can also be used when
  - Process waits for I/O
  - Process waits for another process
- Example of cooperation among processes: UNIX pipes

# Two Approaches To Process Coordination

- Use hardware mechanisms
  - Most useful for multiprocessor systems
  - May rely on *busy waiting*
- Use operating system mechanisms
  - Works well with a single processor
  - No unnecessary execution

Note: we will mention hardware quickly, and focus on operating system functions

# Key Situations That Process Coordination Mechanisms Handle

- Producer / consumer interaction
- Mutual exclusion

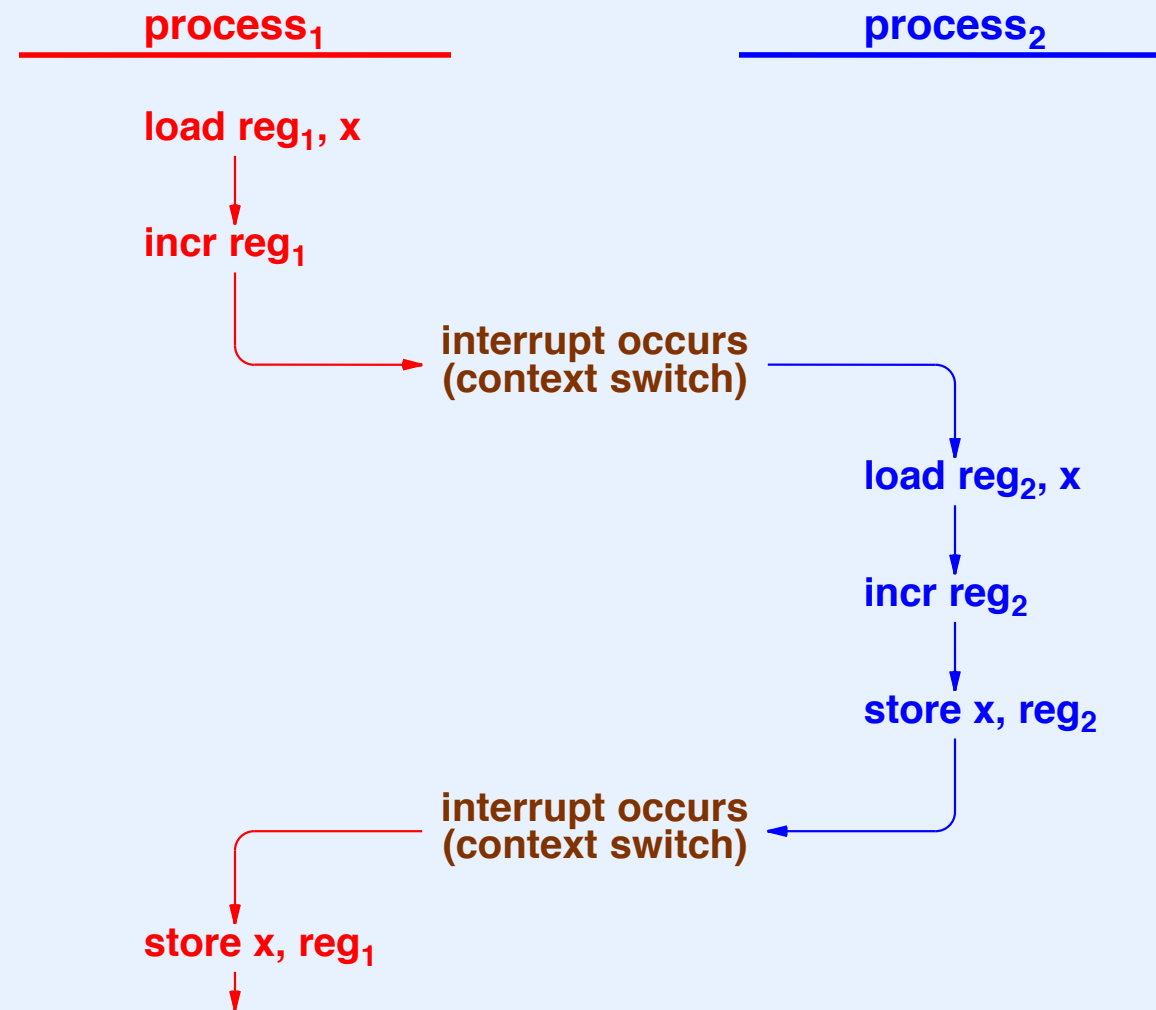
# Producer-Consumer Synchronization

- Typical scenario
  - Shared circular buffer
  - Producing processes deposit items into buffer
  - Consuming processes extract items from buffer
- Must guarantee
  - A producer blocks when buffer full
  - A consumer blocks when buffer empty

# Mutual Exclusion

- Concurrent processes access shared data
- Nonatomic operations can produce unexpected results
- Example: multiple steps used to increment variable  $z$ 
  - Load variable  $z$  into register  $i$
  - Increment register  $i$
  - Store register  $i$  in variable  $z$

# Illustration Of Two Processes Attempting To Increment A Shared Variable Concurrently





# To Prevent Problems

- Ensure that only one process accesses a shared item at any time
- Trick: once a process obtains access, make all other processes wait
- Three solutions
  - Spin lock hardware instructions
  - Disabling all interrupts
  - Semaphores (implemented in software)

# Handling Mutual Exclusion With Spin Locks

- Used in multiprocessors; does not work for single processor
- Atomic hardware operation tests a memory location and changes it
- Known as *test-and-set*
- Also called *spin lock* because it involves *busy waiting*

## Example Of A Spin Lock (x86)

- Instruction performs atomic compare and exchange (*cmpxchg*)
- Spin loop: repeat the following
  - Place a “unlocked” value (e.g, 0) in register *eax*
  - Place an “locked” value (e.g., 1) in register *ebx*
  - Place the address of a memory location in register *ecx* (the lock)
  - Execute the *cmpxchg* instruction
  - Register *eax* will contain the value of the lock before the compare and exchange occurred
  - Continue the spin loop as long as *eax* contains the “locked” value
- To release, assign the “unlocked” value to the lock

# Example Spin Lock Code For X86 (part 1)

```
/* mutex.S - mutex_lock, mutex_unlock */

    .text
    .globl mutex_lock
    .globl mutex_unlock

/*-----
 * mutex_lock(uint32 *lock)  --  Acquire a lock
 *-----
 */
mutex_lock:

    /* Save registers that will be modified */

    pushl    %eax
    pushl    %ebx
    pushl    %ecx
```

## Example Spin Lock Code For X86 (part 2)

```
spinloop:
    movl    $0, %eax          /* Place the "unlocked" value in eax    */
    movl    $1, %ebx          /* Place the "locked" value in ebx     */
    movl    16(%esp), %ecx     /* Place the address of the lock in ecx */

    lock    cmpxchg %ebx, (%ecx) /* Atomic compare-and-exchange:        */
                                /* Compare ebx with memory (%ecx)      */
                                /* if equal                            */
                                /*      load %ebx in memory (%ecx)     */
                                /* else                                */
                                /*      load %ebx in %eax              */

    /* If eax is 1, the mutex was locked, so continue the spin loop */

    cmp     $1, %eax
    je      spinloop

    /* We hold the lock now, so pop the saved registers and return */
    popl    %ecx
    popl    %ebx
    popl    %eax
    ret
```

## Example Spin Lock Code For X86 (part 3)

```
/*-----  
 * mutex_unlock (uint32 *lock) - release a lock  
 *-----  
 */  
mutex_unlock:  
  
    /* Save register eax */  
    pushl    %eax  
  
    /* Load the address of lock onto eax */  
    movl     8(%esp), %eax  
  
    /* Store the "unlocked" value in the lock, thereby unlocking it */  
    movl     $0, (%eax)  
  
    /* Restore the saved register and return */  
    popl     %eax  
    ret
```

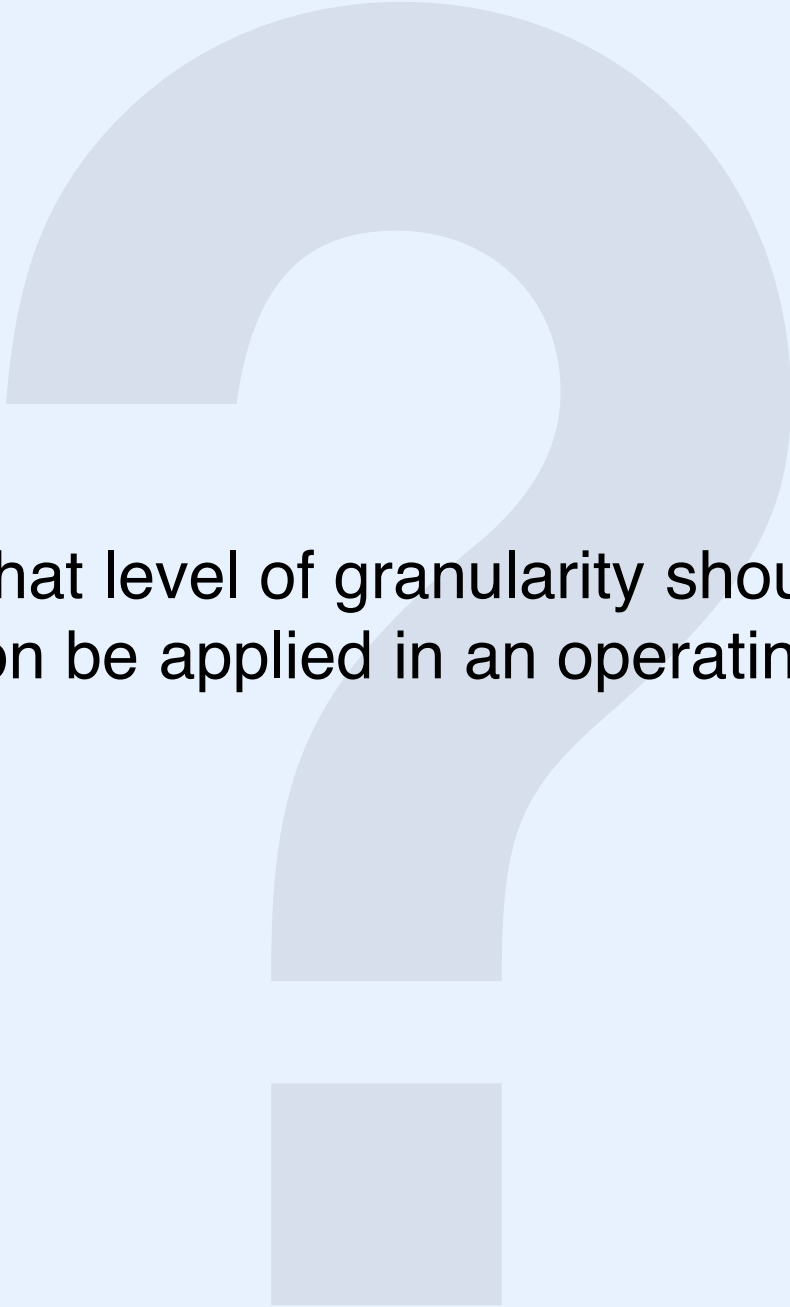
# Handling Mutual Exclusion With Semaphores

- Semaphore allocated for item to be protected
- Known as a *mutex* semaphore
- Applications must be programmed to use the mutex semaphore before accessing shared item
- Operating system guarantees only one process can access the shared item at a given time
- Implementation avoids busy waiting

# Definition Of Critical Section

- Each piece of shared data must be protected from concurrent access
- Programmer inserts mutex operations
  - Before access to shared item
  - After access to shared item
- Protected code known as *critical section*
- Mutex operations can be placed in each function that accesses the shared item





At what level of granularity should mutual exclusion be applied in an operating system?

# Low-Level Mutual Exclusion

- Mutual exclusion needed
  - By application processes
  - Inside operating system
- Mutual exclusion can be guaranteed provided no context switching occurs
- Context changed by
  - Interrupts
  - Calls to *resched*
- Low-level mutual exclusion: mask interrupts and avoid rescheduling

# Interrupt Mask

- Hardware mechanism that controls interrupts
- Internal machine register; may be part of processor status word
- On some hardware, zero value means interrupts can occur; on other hardware, one means interrupts can occur
- OS can
  - Examine current interrupt mask (find out whether interrupts are enabled)
  - Set interrupt mask to prevent interrupts
  - Clear interrupt mask to allow interrupts

# Masking Interrupts

- Important principle:

No operating system function should contain code to explicitly enable interrupts.

- Technique used: given function
  - Saves current interrupt status
  - Disables interrupts
  - Proceeds through critical section
  - Restores interrupt status from saved copy
- Key insight: save / restore allows arbitrary call nesting

# Why Interrupt Masking Is Insufficient

- It works! But...
- Stopping interrupts penalizes all processes when one process executes a critical section
  - Stops all I/O activity
  - Restricts execution to one process for the entire system
- Can interfere with the scheduling invariant (low-priority process can block a high-priority process for which I/O has completed)
- Does not provide a policy that controls which process can access the critical section at a given time

# High-Level Mutual Exclusion

- Idea is to create a facility with the following properties
  - Permit designer to specify multiple critical sections
  - Allow independent control of each critical section
  - Provide an access policy (e.g., FIFO)
- A single mechanism, the *counting semaphore*, suffices

# Counting Semaphore

- Operating system abstraction
- Instance can be created dynamically
- Each instance given unique name
  - Typically an integer
  - Known as a *semaphore ID*
- Instance consists of a tuple (count, set)
  - *Count* is an integer
  - *Set* is a set of processes waiting on the semaphore

# Operations On Semaphores

- *Create* a new semaphore
- *Delete* an existing semaphore
- *Wait* on an existing semaphore
  - Decrements count
  - Adds calling process to set of waiting processes if resulting count is negative
- *Signal* an existing semaphore
  - Increments count
  - Makes a process ready if any are waiting



# Key Uses Of Counting Semaphores

- Two basic paradigms
  - Cooperative mutual exclusion
  - Direct synchronization (e.g., producer-consumer)

# Mutual Exclusion With Semaphores

- Initialize: create a mutex semaphore

```
sid = semcreate(1);
```

- Use: bracket critical sections of code with calls to *wait* and *signal*

```
wait(sid);  
...critical section (use shared resource)...  
signal(sid);
```

- Guarantees only one process can access the critical section at any time (others will be blocked)

# Producer-Consumer Synchronization With Semaphores

- Two semaphores suffice
- Initialize: create producer and consumer semaphores

```
psem = semcreate(buffer-size);  
csem = semcreate(0);
```

- Producer algorithm

```
repeat forever {  
    wait(psem);  
    fill_next_buffer_slot;  
    signal(csem);  
}
```


# Producer-Consumer Synchronization

## (continued)

- Consumer algorithm

```
repeat forever {  
    wait(csem);  
    extract_from_buffer_slot;  
    signal(psem);  
}
```

# Interpretation Of Producer-Consumer Semaphores



- *csem* counts items currently in buffer
- *psem* counts unused slots in buffer

# Semaphore Invariant

# Semaphore Invariant

- Establishes relationship between semaphore concept and implementation

# Semaphore Invariant

- Establishes relationship between semaphore concept and implementation
- Makes code easy to create and understand



# Semaphore Invariant

- Establishes relationship between semaphore concept and implementation
- Makes code easy to create and understand
- Must be re-established after each operation

# Semaphore Invariant

- Establishes relationship between semaphore concept and implementation
- Makes code easy to create and understand
- Must be re-established after each operation
- Surprisingly elegant:

**A nonnegative semaphore count means that the set of processes is empty. A count of negative  $N$  means that the set contains  $N$  waiting processes.**

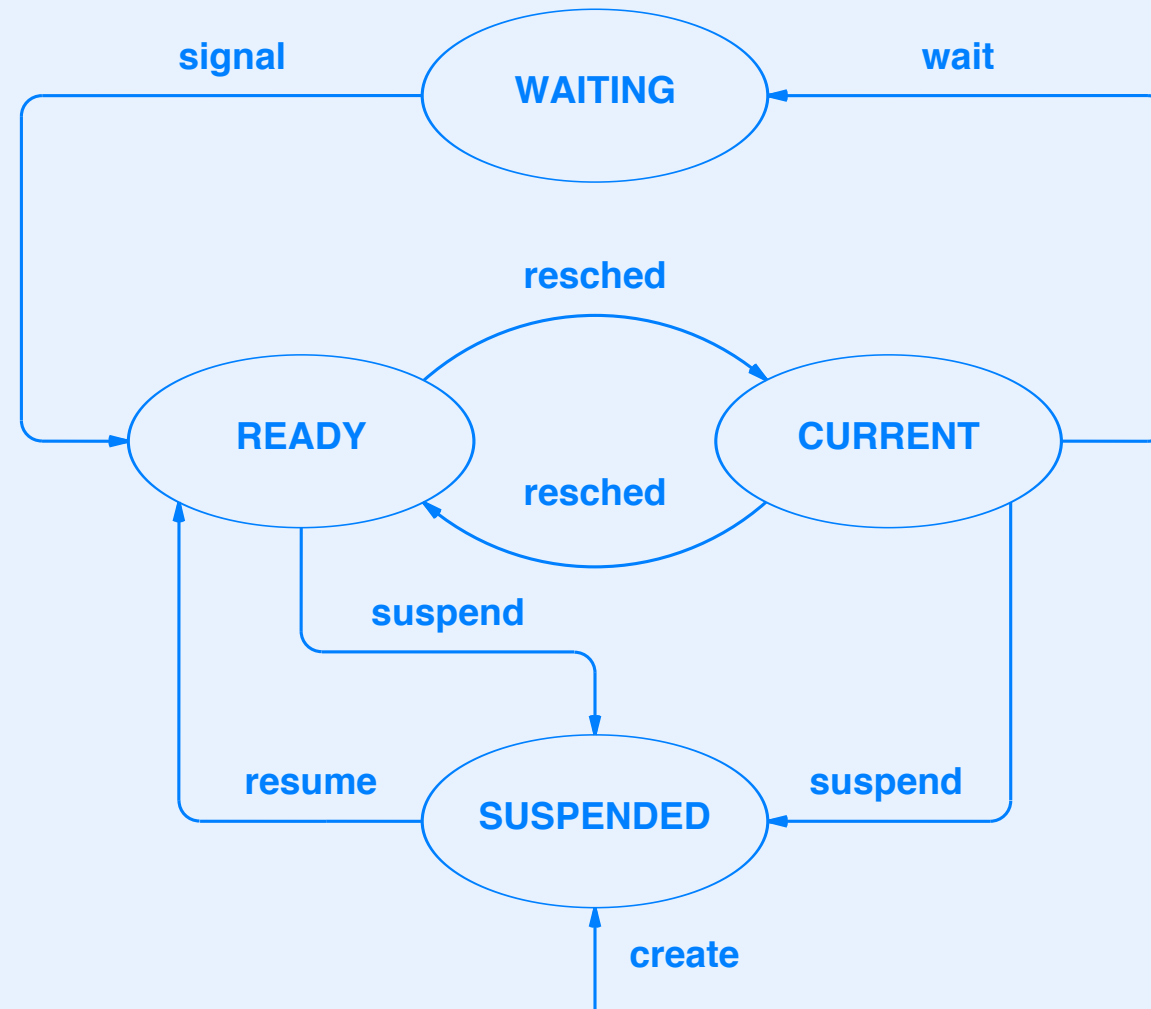
# Counting Semaphores In Xinu

- Stored in an array of semaphore entries
- Each entry
  - Corresponds to one instance (one semaphore)
  - Contains an integer count and pointer to list of processes
- Semaphore ID is index into array
- Policy for management of waiting processes is FIFO

# Process State Used With Semaphores

- When process is waiting on a semaphore, process is not
  - Executing
  - Ready
  - Suspended
- Suspended state is only used by *suspend* and *resume*
- Therefore a new state is needed
- We will use the *WAITING* state for a process blocked a semaphore

# State Transitions With Waiting State



# Semaphore Definitions

```
/* semaphore.h - isbadsem */

#ifndef NSEM
#define NSEM          120      /* Number of semaphores, if not defined */
#endif

/* Semaphore state definitions */

#define S_FREE    0      /* Semaphore table entry is available */
#define S_USED    1      /* Semaphore table entry is in use */

/* Semaphore table entry */
struct sentry {
    byte    sstate;      /* Whether entry is S_FREE or S_USED */
    int32    scount;      /* Count for the semaphore */
    qid16    squeue;      /* Queue of processes that are waiting */
                        /*      on the semaphore */
};

extern struct sentry semtab[];

#define isbadsem(s)      ((int32)(s) < 0 || (s) >= NSEM)
```

# Implementation Of Wait (part 1)

```
/* wait.c - wait */

#include <xinu.h>

/*-----
 * wait - Cause current process to wait on a semaphore
 *-----
 */
syscall wait(
    sid32      sem          /* Semaphore on which to wait */
)
{
    intmask mask;           /* Saved interrupt mask */
    struct procent *prptr;  /* Ptr to process' table entry */
    struct sentry *semptr;  /* Ptr to semaphore table entry */

    mask = disable();
    if (isbadsem(sem)) {
        restore(mask);
        return SYSERR;
    }

    semptr = &semtab[sem];
    if (semptr->sstate == S_FREE) {
        restore(mask);
        return SYSERR;
    }
}
```

# Implementation Of Wait (part 2)

```
if (--(semptr->scount) < 0) {           /* If caller must block */
    prptr = &proctab[currpid];
    prptr->prstate = PR_WAIT;           /* Set state to waiting */
    prptr->prsem = sem;                 /* Record semaphore ID */
    enqueue(currpid,semptr->squeue);    /* Enqueue on semaphore */
    resched();                         /* and reschedule */
}

restore(mask);
return OK;
}
```



# Semaphore Queuing Policy

- Determines which process to select among those waiting
- Needed when *signal* called
- Examples
  - First-Come-First-Served (FCFS or FIFO)
  - Process priority
  - Random

# Question

# Question

- The goal is “fairness”

# Question

- The goal is “fairness”
- Which semaphore queuing policy implements goal best?

# Question

- The goal is “fairness”
- Which semaphore queuing policy implements goal best?
- In other words, how should we interpret fairness?

# Question

- The goal is “fairness”
- Which semaphore queuing policy implements goal best?
- In other words, how should we interpret fairness?
- Semaphore policy can interact with scheduling policy

# Question

- The goal is “fairness”
- Which semaphore queuing policy implements goal best?
- In other words, how should we interpret fairness?
- Semaphore policy can interact with scheduling policy
  - Should a low-priority process be allowed to access a resource if a high-priority process is also waiting?
  - Should a low-priority process be blocked forever if high-priority processes use a resource?

# Choosing A Semaphore Queueing Policy

- Difficult
- No single best answer
  - Fairness not easy to define
  - Scheduling and coordination interact in subtle ways
  - May affect other OS policies
- Interactions of heuristic policies may produce unexpected results



# Semaphore Queuing Policy In Xinu

- First-come-first-serve
- Straightforward to implement
- Extremely efficient
- Works well for traditional uses of semaphores
- Potential problem: low-priority process can access a resource while a high-priority process remains blocked

# Implementation Of FIFO Semaphore Policy

- Each semaphore uses a list to manage waiting processes
- List is run as a queue: insertions at one end and deletions at the other
- Example implementation follows

# Implementation Of Signal (part 1)

```
/* signal.c - signal */

#include <xinu.h>

/*-----
 * signal - Signal a semaphore, releasing a process if one is waiting
 *-----
 */
syscall signal(
    sid32      sem      /* ID of semaphore to signal */
)
{
    intmask mask;          /* Saved interrupt mask */
    struct sentry *semptr; /* Ptr to semaphore table entry */

    mask = disable();
    if (isbadsem(sem)) {
        restore(mask);
        return SYSERR;
    }
    semptr = &semtab[sem];
    if (semptr->sstate == S_FREE) {
        restore(mask);
        return SYSERR;
    }
}
```

# Implementation Of Signal (part 2)

```
if ((semptr->scount++) < 0) {    /* Release a waiting process */  
    ready(dequeue(semptr->squeue));  
}  
restore(mask);  
return OK;  
}
```

# Semaphore Allocation

- Static
  - Semaphores are defined at compile time
  - More efficient, but less powerful
- Dynamic
  - Semaphores are created at runtime
  - More flexible
- Xinu supports dynamic allocation

# Xinu Semcreate (part 1)

```
/* semcreate.c - semcreate, newsem */

#include <xinu.h>

local  sid32  newsem(void);

/*-----
 * semcreate - Create a new semaphore and return the ID to the caller
 *-----
 */
sid32  semcreate(
        int32          count          /* Initial semaphore count          */
)
{
    intmask mask;                    /* Saved interrupt mask          */
    sid32  sem;                      /* Semaphore ID to return        */

    mask = disable();

    if (count < 0 || ((sem=newsem())==SYSERR)) {
        restore(mask);
        return SYSERR;
    }
    semtab[sem].scount = count;      /* Initialize table entry        */

    restore(mask);
    return sem;
}
```

## Xinu Semcreate (part 2)

```
/*-----  
 * newsem - Allocate an unused semaphore and return its index  
 *-----  
 */  
local sid32 newsem(void)  
{  
    static sid32 nextsem = 0; /* Next semaphore index to try */  
    sid32 sem; /* Semaphore ID to return */  
    int32 i; /* Iterate through # entries */  
  
    for (i=0 ; i<NSEM ; i++) {  
        sem = nextsem++;  
        if (nextsem >= NSEM)  
            nextsem = 0;  
        if (semtab[sem].sstate == S_FREE) {  
            semtab[sem].sstate = S_USED;  
            return sem;  
        }  
    }  
    return SYSERR;  
}
```

# Semaphore Deletion

- Wrinkle: one or more processes may be waiting when semaphore is deleted
- Must choose a disposition for each
- Xinu policy: make process ready



# Xinu Semdelete (part 1)

```
/* semdelete.c - semdelete */

#include <xinu.h>

/*-----
 * semdelete - Delete a semaphore by releasing its table entry
 *-----
 */
syscall semdelete(
    sid32      sem      /* ID of semaphore to delete */
)
{
    intmask mask;        /* Saved interrupt mask */
    struct sentry *semptr; /* Ptr to semaphore table entry */

    mask = disable();
    if (isbadsem(sem)) {
        restore(mask);
        return SYSERR;
    }

    semptr = &semtab[sem];
    if (semptr->sstate == S_FREE) {
        restore(mask);
        return SYSERR;
    }
    semptr->sstate = S_FREE;
```

## Xinu Semdelete (part 2)

```
    resched_cntl(DEFER_START);
    while (semptr->scount++ < 0) { /* Free all waiting processes */
        ready(getfirst(semptr->squeue));
    }
    resched_cntl(DEFER_STOP);
    restore(mask);
    return OK;
}
```




Do you understand semaphores?

# **Module VI**

## **Inter-Process Communication**

# Location Of Inter-process Communication In The Hierarchy



# Inter-process Communication

- Used for
  - Exchange of (nonshared) data
  - Process coordination
- General technique: *message passing*

# Two Approaches To Message Passing

- Approach #1
  - Message passing is one of many services
  - Messages are separate from I/O and process synchronization services
  - Messages implemented using lower-level mechanisms, such as semaphores
- Approach #2
  - The entire operating system is *message-based*
  - Messages, not function calls, provide the fundamental building block
  - Messages, not semaphores, used for process synchronization

# Design Of A Message Passing Facility



# Design Of A Message Passing Facility

- To understand the issues, we will begin with a trivial message passing facility

# Design Of A Message Passing Facility

- To understand the issues, we will begin with a trivial message passing facility
- We want to allow a process to send a message directly to another process

# Design Of A Message Passing Facility

- To understand the issues, we will begin with a trivial message passing facility
- We want to allow a process to send a message directly to another process
- In principle,

# Design Of A Message Passing Facility

- To understand the issues, we will begin with a trivial message passing facility
- We want to allow a process to send a message directly to another process
- In principle, the design should be straightforward

# Design Of A Message Passing Facility

- To understand the issues, we will begin with a trivial message passing facility
- We want to allow a process to send a message directly to another process
- In principle, the design should be straightforward
- In practice,

# Design Of A Message Passing Facility

- To understand the issues, we will begin with a trivial message passing facility
- We want to allow a process to send a message directly to another process
- In principle, the design should be straightforward
- In practice, many design decisions arise

# Message Passing Design Decisions

- Are messages fixed or variable size?
- What is the maximum message size?
- How many messages can be outstanding at a given time?
- Where are messages stored?
- How is a recipient specified?
- Does a receiver know the sender's identity?
- Are replies supported?
- Is the interface synchronous or asynchronous?

# Synchronous vs. Asynchronous Interface

- Synchronous interface
  - Blocks until the operation is performed
  - Easy to understand / program
  - Extra processes can be used to obtain asynchrony



# Synchronous vs. Asynchronous Interface (continued)

- Asynchronous interface
  - Process starts an operation
  - Initiating process continues execution
  - Notification arrives when operation completes
    - \* Happens at any time
    - \* May entail abnormal control flow (e.g., software interrupt or “callback” mechanism)
  - Polling can be used to determine status

# Why Is A Message Passing Facility So Difficult To Design?

- Interacts with
  - Process coordination subsystem
  - Memory management subsystem
- Affects user's perception of system

# Example Inter-process Message Passing Design

- Simple, low-level mechanism
- Direct process-to-process communication
- One-word messages
- Message stored with receiver
- One-message buffer
- Synchronous, buffered reception
- Asynchronous transmission and “reset” operation

# Example Inter-process Message Passing Design (continued)

- Three functions

```
send(msg, pid);
```

```
msg = receive();
```

```
msg = recvclr();
```

- Message stored in *receiver's* process table entry
- *Send* transmits message to specified process
- *Receive* blocks until a message arrives
- *Recvclr* removes existing message, if one has arrived, but does not block

# Example Inter-process Message Passing Design (continued)

- First-message semantics

- First message sent to a process is stored until it has been received
- Subsequent attempts to send fail

- Idiom


```
recvclr(); /* prepare to receive a message */  
... /* allow other processes to send messages */  
msg = receive();
```

- Above code returns first message that was sent, even if a high priority process attempts to send later

# Process State For Message Reception

- While receiving a message, a process is not
  - Executing
  - Ready
  - Suspended
  - Waiting on a semaphore
- Therefore, a new state is needed for message passing
- Named *RECEIVING*
- Entered when *receive* called

# State Transitions With Message Passing



# Xinu Code For Message Reception

```
/* receive.c - receive */

#include <xinu.h>

/*-----
 * receive - Wait for a message and return the message to the caller
 *-----
 */
umsg32 receive(void)
{
    intmask mask;                /* Saved interrupt mask */
    struct procent *prptr;        /* Ptr to process' table entry */
    umsg32 msg;                   /* Message to return */

    mask = disable();
    prptr = &proctab[currpid];
    if (prptr->prhasmsg == FALSE) {
        prptr->prstate = PR_RECV;
        resched();                /* Block until message arrives */
    }
    msg = prptr->prmsg;            /* Retrieve message */
    prptr->prhasmsg = FALSE;       /* Reset message flag */
    restore(mask);
    return msg;
}
```



# Xinu Code For Message Transmission (part 1)

```
/* send.c - send */

#include <xinu.h>

/*-----
 * send - Pass a message to a process and start recipient if waiting
 *-----
 */
syscall send(
    pid32      pid,      /* ID of recipient process */
    umsg32     msg       /* Contents of message      */
)
{
    intmask mask;        /* Saved interrupt mask     */
    struct procent *prptr; /* Ptr to process' table entry */

    mask = disable();
    if (isbadpid(pid)) {
        restore(mask);
        return SYSERR;
    }

    prptr = &proctab[pid];
    if ((prptr->prstate == PR_FREE) || prptr->prhasmsg) {
        restore(mask);
        return SYSERR;
    }
}
```

## Xinu Code For Message Transmission (part 2)

```
prptr->prmsg = msg;           /* Deliver message */
prptr->prhasmsg = TRUE;        /* Indicate message is waiting */

/* If recipient waiting or in timed-wait make it ready */

if (prptr->prstate == PR_RECV) {
    ready(pid);
} else if (prptr->prstate == PR_RECTIM) {
    unsleep(pid);
    ready(pid);
}
restore(mask);                /* Restore interrupts */
return OK;
}
```

- Note: we will discuss receive-with-timeout later

# Xinu Code For Clearing Messages

```
/* recvclr.c - recvclr */

#include <xinu.h>

/*-----
 *  recvclr  -  Clear incoming message, and return message if one waiting
 *-----
 */
umsg32  recvclr(void)
{
    intmask mask;                /* Saved interrupt mask */
    struct procent *prptr;       /* Ptr to process' table entry */
    umsg32  msg;                 /* Message to return */

    mask = disable();
    prptr = &proctab[currpid];
    if (prptr->prhasmsg == TRUE) {
        msg = prptr->prmsg;      /* Retrieve message */
        prptr->prhasmsg = FALSE; /* Reset message flag */
    } else {
        msg = OK;
    }
    restore(mask);
    return msg;
}
```

# Summary

- Inter-process communication
  - Implemented by message passing
  - Can be synchronous or asynchronous
- Synchronous interface is the simplest
- Xinu uses synchronous reception and asynchronous transmission



**Questions?**