SOLUTIONS TO CS 354 MIDTERM, SPRING 2013 (PARK)

P1(a) 12 pts

Disabling interrupts guarantees that a system call runs uninterrupted
which ensures integrity of shared kernel data structures through mutual
exclusion.
4 pts

create() has to update the kernel's process table. Without interrupt
disabling, the process table may get corrupted if two processes executing
create() try to access it concurrently.
4 pts

Interrupt disabling is a drastic action that can prevent other important
events, such as handling packets arriving at a network interface card
and clock interrupts, from being processed by the kernel in a timely manner.
4 pts

P1(b) 12 pts

Three hardware features: privileged/non-privileged instruction set,
kernel/user mode, memory protection.
6 pts

System calls (i) change a process's status from user mode to kernel mode,
(ii) branch to a specific kernel function to handle the user request
in kernel space. The kernel function may include privileged instructions
for accessing shared hardware resources which can now be executed without
triggering a fault. When a system call returns the process is switched
back to user mode.
6 pts

P1(c) 12 pts

Four pieces: the caller's frame pointer (ebp), flags register bits, 8
general purpose registers, stack pointer.
6 pts

The stack pointer of a context-switched out process needs to be easily
accessible when context-switching in the process at a later time. The process
table serves this purpose well.
3 pts

Accessing the process table entails several instructions for memory
indirection. Using hardware supplied instructions (push/pop variety for
stack manipulation) is more efficient.
3 pts

P2(a) 16 pts

Processes that make blocking system calls and give up CPU voluntarily
are treated as IO-intensive and their priority is increased while their
time slice is decreased. Processes that deplete their time slice and
are preempted by the clock interrupt handler are treated as CPU-intensive
processes and their priority is decreased and their time slice increased.
6 pts

The rationale is to give I/O-intensive processes preference over
CPU-intensive processes when they are in ready state, since by their
nature, they consume less CPU cycles. A small time slice is given to
I/O-intensive processes, just in case a CPU-intensive process was
misclassified.
6 pts

A potential issue is that for real-world processes that are a mix --
exhibit both CPU- and I/O-intensive behavior over time -- how they fare
relative to CPU- and I/O-intensive processes is less clear.

2 pts

Solaris monitors how long a ready process has not received CPU cycles,
and if this time period exceeds a threshold, its priority is increased.
2 pts

P2(b) 16 pts

Pro: Thread creation and management does not entail system calls which
tend to be slow.
6 pts

Cons: (i) User space threads cannot make blocking system calls (e.g.,
for I/O) since a single thread blocking ends up blocking all user space
threads. (ii) Multiple CPUs (or cores) cannot be utilized.
8 pts

Due to (ii), multithreading with kernel support is more suited for multicore
processors.
2 pts

P3(a) 16 pts

Xinu's default scheduler requires linear over head (i.e., O(n) where n
is the number of processes) since inserting a process into a priority
queue, in the worst case, requires traversing the whole list.
6 pts

With k priority levels, where k is treated as a constant: (i) Dequeue
requires finding the highest priority level at which there are one or more
ready processes enqueued. At worst, this requires k comparisons. After
finding the highest non-empty priority level, we can pick the first
process (constant cost) since all processes at the same level are serviced
round-robin. (ii) Enqueue requires going to a process's priority level
(constant since multi-level feedback queue is an array of linked lists
indexed by priority), which is constant, then inserting at the end of
the list (also constant) by round-robin property.
10 pts

P3(b) 16 pts

tset does not disable interrupts, thus important events such as packet
arrivals can continue to be processed by a kernel's interrupt handling
routines.
6 pts

tset wastes CPU cycles by waiting on tset to return false in an infinite
loop (busy waiting). A process that uses semaphores does not busy wait
but gets blocked (and context-switched out) by the kernel which avoids
wasting CPU cycles.
6 pts

Yes, semaphore calls wait() and signal(), internally, must disable
interrupts to ensure that their code is run atomically. Since the codes
of wait() and signal() are relatively short, the disruption to kernel
due to interrupt disabling is kept small.
4 pts

Bonus 10 pts

Without a kernel stack, two processes A and B that share memory where
their run-time stacks reside can collude and run their own code in kernel
mode. For example, A makes a blocking system call, say, read(). read()
may entail internal kernel function calls which are managed (push/pop)
using A's user space run-time stack. read() blocks and A is context
switched out while in kernel mode. Process B is scheduled next, B has
access to A's run-time stack, overwrites one of the return addresses
(EIP) of read()'s callees so that it points to B's own code. When A

unblocks and is context switched in again, returning from read()'s nested
function calls will cause a jump to B's code while A is still in kernel
mode.